

기말 프로젝트

CNN 에 기반한 image classification 알고리즘 개발

201711809 박수빈

45000개의 이미지를 numpy array로 변환



exterior



food



interior



sample

Images폴더에 들어있던 음식, 실내, 실외사진을 폴더를 각각 만들어서 분류했습니다.

그리고 모델 학습 후 예측할 sample사진 10개도 폴더를 만들어서 저장했습니다.

Cv2의 imread함수를 사용해서 이미지파일을 읽어 리스트에 저장했습니다.

Y 배열에는 음식0, 실내1, 실외2로 저장했습니다.

모든 파일을 읽어서 리스트에 추가했으면 np.array로 형식을 변환해줍니다.

```
def train_mnist_model():
    X, y = [], []
    for i in range(20000):
        X.append(cv2.imread('images/food/food%d.jpg'%(i+1)))
        y.append(0)
    print('food insert finished')

    for i in range(15000):
        X.append(cv2.imread('images/interior/interior%d.jpg'%(i+1)))
        y.append(1)
    print('interior insert finished')

    for i in range(10000):
        X.append(cv2.imread('images/exterior/exterior%d.jpg'%(i+1)))
        y.append(2)
    print('exterior insert finished')

    X = np.array(X)
    y = np.array(y)
```

train/test 데이터 구분 및 셔플

```
X_train, X_test, y_train, y_test = train_test_split(X, y, test_size=0.33, random_state=42)

y_train = to_categorical(y_train)
y_test = to_categorical(y_test)
```

X와 y는 총 45000개의 data가 들어있고 학습과 테스트모델을 분류하고 데이터가 골고루 들어가기 위해서 train_test_split함수를 사용해 shuffle했습니다.

나뉘진 y_test와 y_train 배열은 0, 1, 2로 이루어져있는데 원핫인코딩하기 위해 to_categorical함수를 사용했습니다.

train data 및 validation data의 loss/accuracy 분석

```
def plot_loss_curve(history):  
  
    plt.figure(figsize=(15, 10))  
  
    plt.plot(history['loss'])  
    plt.plot(history['val_loss'])  
    plt.title('model loss')  
    plt.ylabel('loss')  
    plt.xlabel('epoch')  
    plt.legend(['train', 'test'], loc='upper right')  
    plt.show()  
  
def plot_accuracy_curve(history):  
  
    plt.figure(figsize=(15, 10))  
  
    plt.plot(history['accuracy'])  
    plt.plot(history['val_accuracy'])  
    plt.title('model accuracy')  
    plt.ylabel('accuracy')  
    plt.xlabel('epoch')  
    plt.legend(['train', 'test'], loc='upper right')  
    plt.show()
```



plot_loss_curve는 loss와 val_loss를 epoch가 지나면서 변화를 그래프로 나타내는 함수입니다.

plot_accuracy_curve는 accuracy 와 val_accuracy를 epoch가 지나면서 변화를 그래프로 나타내는 함수입니다.

train data 및 validation data의 loss/accuracy분석

```
Epoch 1/10
302/302 [=====] - 351s 1s/step - loss: 277.4598 - accuracy: 0.3888 - val_loss: 1.0864 - val_accuracy: 0.4754
Epoch 2/10
302/302 [=====] - 394s 1s/step - loss: 0.9163 - accuracy: 0.5688 - val_loss: 1.0391 - val_accuracy: 0.5044
Epoch 3/10
302/302 [=====] - 352s 1s/step - loss: 0.6052 - accuracy: 0.7364 - val_loss: 1.1677 - val_accuracy: 0.5320
Epoch 4/10
302/302 [=====] - 358s 1s/step - loss: 0.3732 - accuracy: 0.8501 - val_loss: 1.4186 - val_accuracy: 0.5205
Epoch 5/10
302/302 [=====] - 337s 1s/step - loss: 0.2896 - accuracy: 0.8896 - val_loss: 1.7894 - val_accuracy: 0.5165
Epoch 6/10
302/302 [=====] - 332s 1s/step - loss: 0.3277 - accuracy: 0.9082 - val_loss: 1.5286 - val_accuracy: 0.5239
Epoch 7/10
302/302 [=====] - 326s 1s/step - loss: 0.2728 - accuracy: 0.9157 - val_loss: 2.3479 - val_accuracy: 0.5394
Epoch 8/10
302/302 [=====] - 346s 1s/step - loss: 0.2481 - accuracy: 0.9238 - val_loss: 2.0734 - val_accuracy: 0.4909
Epoch 9/10
302/302 [=====] - 345s 1s/step - loss: 0.2454 - accuracy: 0.9230 - val_loss: 2.6707 - val_accuracy: 0.4923
Epoch 10/10
302/302 [=====] - 346s 1s/step - loss: 0.2731 - accuracy: 0.9146 - val_loss: 2.5483 - val_accuracy: 0.4983
```

train data 및 validation data의 loss/accuracy 분석



Model 구성

```
model = Sequential([
    Input(shape=(300,300,3), name='input_layer'),
    Conv2D(32, kernel_size=3, activation='relu', name='conv_layer1'),
    MaxPooling2D(pool_size=2),
    Conv2D(64, kernel_size=3, activation='relu', name='conv_layer2'),
    MaxPooling2D(pool_size=2),
    Dropout(0.5),
    Flatten(),
    Dense(64, activation='relu', name='output_layer1'),
    Dense(3, activation='softmax', name='output_layer2')
])
```

Layer을 추가하면 파라미터 개수가 너무 많아져서 시간이 많이걸려서 layer의 개수를 몇개할지 고민하다가 2개로 결정했습니다.

Conv2d 3*3필터를 2개 사용했고 활성화함수로 relu함수를 사용했고 사진을 너무 자세히 보지 않고 크게보게하기 위해 maxpooling2d필터를 사용했습니다.

오버피팅을 막기 위해서 dropout함수를 사용했습니다.

Model 구성

Layer (type)	Output Shape	Param #
conv_layer1 (Conv2D)	(None, 298, 298, 32)	896
max_pooling2d (MaxPooling2D)	(None, 149, 149, 32)	0
conv_layer2 (Conv2D)	(None, 147, 147, 64)	18496
max_pooling2d_1 (MaxPooling2D)	(None, 73, 73, 64)	0
dropout (Dropout)	(None, 73, 73, 64)	0
flatten (Flatten)	(None, 341056)	0
output_layer1 (Dense)	(None, 64)	21827648
output_layer2 (Dense)	(None, 3)	195
Total params: 21,847,235		
Trainable params: 21,847,235		
Non-trainable params: 0		


Prediction 시연



sample1




sample2



sample3




sample4



sample5



sample6



sample7



sample8



sample9



sample10

Sample사진들은 제가 직접 찍은 사진들로 음식사진4개 실내사진3개 실외사진 3개로 총 10개로 구성했습니다.

Prediction 시연



Prediction 시연



성능을 높이기 위해 노력한 방법들

모델의 layer구성도 다르게 해보고 epoch수와 batch사이즈도 변경하면서 테스트했는데 정확도는 거의 50퍼로 비슷하게 나왔습니다.

사진이 45000이기 때문에 시간이 너무 오래 걸려서 4500장으로 추려서 테스트를 많이 해봤습니다.

```
Model: "sequential"
Layer (type)                Output Shape                Param #
-----
conv_layer1 (Conv2D)        (None, 298, 298, 32)        896
dropout (Dropout)           (None, 298, 298, 32)        0
max_pooling2d (MaxPooling2D) (None, 149, 149, 32)        0
flatten (Flatten)           (None, 710432)              0
output_layer (Dense)        (None, 3)                   2131299
-----
Total params: 2,132,195
Trainable params: 2,132,195
Non-trainable params: 0

2020-12-17 11:19:29.268000: I tensorflow/compiler/mlir/mlir_graph_optimization_pass.cc:116] None of the MLIR optimization passes are enabled (registered 2)
Epoch 1/10
302/302 [=====] - 277s 917ms/step - loss: 2160.0946 - accuracy: 0.4766 - val_loss: 9.3333 - val_accuracy: 0.5199
Epoch 2/10
302/302 [=====] - 264s 874ms/step - loss: 2.4004 - accuracy: 0.8522 - val_loss: 7.9187 - val_accuracy: 0.5859
Epoch 3/10
302/302 [=====] - 267s 884ms/step - loss: 0.4653 - accuracy: 0.9460 - val_loss: 7.9620 - val_accuracy: 0.5879
Epoch 4/10
302/302 [=====] - 269s 890ms/step - loss: 0.1566 - accuracy: 0.9790 - val_loss: 9.3590 - val_accuracy: 0.5906
Epoch 5/10
302/302 [=====] - 261s 864ms/step - loss: 0.2129 - accuracy: 0.9780 - val_loss: 8.7897 - val_accuracy: 0.5690
Epoch 6/10
302/302 [=====] - 246s 814ms/step - loss: 0.1405 - accuracy: 0.9796 - val_loss: 10.2960 - val_accuracy: 0.5859
Epoch 7/10
302/302 [=====] - 235s 778ms/step - loss: 0.0562 - accuracy: 0.9880 - val_loss: 8.9190 - val_accuracy: 0.5778
Epoch 8/10
302/302 [=====] - 236s 780ms/step - loss: 0.1313 - accuracy: 0.9877 - val_loss: 8.6333 - val_accuracy: 0.5609
Epoch 9/10
302/302 [=====] - 299s 991ms/step - loss: 0.1966 - accuracy: 0.9799 - val_loss: 11.4515 - val_accuracy: 0.5623
Epoch 10/10
302/302 [=====] - 235s 778ms/step - loss: 0.5420 - accuracy: 0.9574 - val_loss: 15.3874 - val_accuracy: 0.5805
```

```
2020-12-17 05:17:25.15350: I tensorflow/compiler/mlir/mlir_graph_optimization_pass.cc:116] None of the MLIR optimization passes are enabled (registered 2)
Model: "sequential"
Layer (type)                Output Shape                Param #
-----
conv_layer1 (Conv2D)        (None, 298, 298, 32)        896
max_pooling2d (MaxPooling2D) (None, 149, 149, 32)        0
flatten (Flatten)           (None, 710432)              0
output_layer (Dense)        (None, 3)                   2131299
-----
Total params: 2,132,195
Trainable params: 2,132,195
Non-trainable params: 0

2020-12-17 05:17:25.857508: I tensorflow/compiler/mlir/mlir_graph_optimization_pass.cc:116] None of the MLIR optimization passes are enabled (registered 2)
Epoch 1/3
3015/3015 [=====] - 1403s 465ms/step - loss: 283.8305 - accuracy: 0.5232 - val_loss: 1.0808 - val_accuracy: 0.5244
Epoch 2/3
3015/3015 [=====] - 1443s 479ms/step - loss: 0.5292 - accuracy: 0.8020 - val_loss: 1.6497 - val_accuracy: 0.5370
Epoch 3/3
3015/3015 [=====] - 1460s 484ms/step - loss: 0.4652 - accuracy: 0.8562 - val_loss: 1.7208 - val_accuracy: 0.4914
{'loss': [48.97663497924805, 0.5803033113479614, 0.5584049224853516], 'accuracy': [0.5273631811141968, 0.7763184309005737, 0.8064345121383667], 'val_loss': [1.0808273553848267, 1.649685025215149, 1.7208101749420166], 'val_accuracy': [0.5243771076202393, 0.5370370149612427, 0.4914478063583374]}
train loss= 0.5584049224853516
validation loss= 1.7208101749420166
2020-12-17 06:31:44.278592: W tensorflow/python/util/util.cc:348] Sets are not currently considered sequences, but this may change in the future, so consider avoiding using them.
```

감사합니다